


Loading and saving PyTorch model weights



(1) Learn to save and load a pretrained model
`torch.save(model.state_dict(), "model.pth")`

Dictionary mapping each layer to its parameter ↓ file name

* `model.load_state_dict(torch.load("model.pth"))`

(2) Saving optimizer state is also recommended
`optimizer.state_dict()`

↓ hyperparameters ↓ historical data, such as past gradients

For AdamW optimizer:

Bop_avg: Exponentially moving average of the gradient used to smooth the gradients)

Bop_avg_sq: Exponentially moving average of the squared gradient used to scale the lr for individual parameters)
 other running estimates needed by optimizer to perform updates correctly